

Tidy Tuesday Freedom Dataset

M. Zhang

2022-03-01 (Update 2022-03-09)

Contents

Introduction	1
Data Details	1
Data Loading	2
Data Dictionary	2
Data Exploration and wrangling	3
Examining the raw data	3
Checking for NA values	5
Recoding	5
Grouping candidates	5
Visual Exploration	5
Developing Questions	9
Data Visualizations	12
Trends for top and worst 5, and choice	13
Overall trends in freedom over time	16
Overall trends in freedom over time by region	18
Additional Analysis joined with other data sets	21
Worldbank Data	21
Conclusion	25

Introduction

We will be working with the `freedom.csv` dataset via the Tidy Tuesday repository. This analysis will not be styled as a “final report”, but rather as a quick walk through of some data analysis and wrangling that took place while exploring the data.

The MSDSO Discord group for the University of Texas Masters in Data Science Online program will be doing weekly explorations of TidyTuesday as an exercise for improving their data science skill sets in a collaborative environment.

For this data set, the analysis was first done independently, then augmented by following along with David Robinson’s livestream video (<https://youtu.be/VOzUHk3aaBw>) for additional learning.

Data Details

This dataset is pulled from the Tidy Tuesday Repository:

Thomas Mock (2022). Tidy Tuesday: A weekly data project aimed at the R ecosystem. <https://github.com/rfordatascience/tidytuesday>.

The original data is from Freedom House and the United Nations via Arthur Cheib.

Freedom House is a nonpartisan organization focused on producing research and reports on themes and trends related to democracy, political rights, and civil liberties. This data set **freedom** contains information in regards to various country's Civil Liberty CL and Political Rights PR index scores, as well as their Least Developed Country **is_ldc** indicator.

Data Loading

```
#freedom <- readr::read_csv('https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2022/2022-02-22/freedom.csv')

freedom <- tt_load("2022-02-22")$freedom %>%
  janitor::clean_names() %>%
  rename(civil_liberties = cl,
         political_rights = pr) %>%
  mutate(country_code = countrycode(country, origin = "country.name", destination = "iso2c"))
```

Downloading file 1 of 1: `freedom.csv`

freedom

```
# A tibble: 4,979 x 9
  country year civil_liberties political_rights status region_code region_name
  <chr>   <dbl>         <dbl>         <dbl> <chr>         <dbl> <chr>
1 Afghan~ 1995             7             7 NF           142 Asia
2 Afghan~ 1996             7             7 NF           142 Asia
3 Afghan~ 1997             7             7 NF           142 Asia
4 Afghan~ 1998             7             7 NF           142 Asia
5 Afghan~ 1999             7             7 NF           142 Asia
6 Afghan~ 2000             7             7 NF           142 Asia
7 Afghan~ 2001             7             7 NF           142 Asia
8 Afghan~ 2002             6             6 NF           142 Asia
9 Afghan~ 2003             6             6 NF           142 Asia
10 Afghan~ 2004             6             5 NF           142 Asia
# ... with 4,969 more rows, and 2 more variables: is_ldc <dbl>,
#   country_code <chr>
```

Data Dictionary

The following tables contain information in regards to the columns available.

freedom.csv

variable	class	description
country	character	Country Name
year	double	Year
CL	double	Civil Liberties
PR	double	Political rights
Status	character	Status (Free F, Not Free NF, Partially Free PF)
Region_Code	double	UN Region code
region_name	character	UN Region Name
is_ldc	double	Is a least developed country (binary 0/1)

The definition for “Least Developed Country” is pulled from the United Nations. A country qualifies for LDC if it meets the criteria for Income, Human Assets, and Economic and Environmental Vulnerability. An

important requirement for inclusion is that the country must agree to the classification to be added to the list.

The Civil Liberties and Political Rights score is generated by FreedomHouse, using a methodology inspired by the Universal Declaration of Human Rights which was adopted by the UN General Assembly in 1948, The Civil Liberties score is a combination of 15 separate indicators, and Political Rights score is a combination of 10 separate indicators. Each of these indicators scale from 0 to 4, with 4 representing the greatest amount of freedom. These scores are summarized into indexes for Civil Liberties `civil_liberties` and Political Rights `political_rights` on a scale from 1 to 7, with 1 representing the greatest freedom.

Status buckets the combined CL and PR scores into 3 general categories: Free, Partially Free, and Not Free. Additional details are available [here](#).

Data Exploration and wrangling

Examining the raw data

We will initially take a precursor inspection of the data, utilizing `summary` for numerical information and `table` for categorical information.

```
glimpse(freedom)
```

```
Rows: 4,979
Columns: 9
$ country      <chr> "Afghanistan", "Afghanistan", "Afghanistan", "Afghani~
$ year         <dbl> 1995, 1996, 1997, 1998, 1999, 2000, 2001, 2002, 2003,~
$ civil_liberties <dbl> 7, 7, 7, 7, 7, 7, 7, 7, 6, 6, 6, 5, 5, 5, 6, 6, 6, 6,~
$ political_rights <dbl> 7, 7, 7, 7, 7, 7, 7, 7, 6, 6, 5, 5, 5, 5, 5, 6, 6, 6, 6,~
$ status       <chr> "NF", "NF", "NF", "NF", "NF", "NF", "NF", "NF", "NF",~
$ region_code  <dbl> 142, 142, 142, 142, 142, 142, 142, 142, 142, 142, 142, 142~
$ region_name  <chr> "Asia", "Asia", "Asia", "Asia", "Asia", "Asia", "Asia~
$ is_ldc       <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,~
$ country_code <chr> "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF",~
```

```
names(freedom)
```

```
[1] "country"      "year"          "civil_liberties" "political_rights"
[5] "status"       "region_code"   "region_name"    "is_ldc"
[9] "country_code"
```

```
freedom %>%
  select(civil_liberties, political_rights, is_ldc) %>%
  summary(freedom)
```

civil_liberties	political_rights	is_ldc
Min. :1.000	Min. :1.000	Min. :0.0000
1st Qu.:2.000	1st Qu.:1.000	1st Qu.:0.0000
Median :3.000	Median :3.000	Median :0.0000
Mean :3.369	Mean :3.411	Mean :0.2362
3rd Qu.:5.000	3rd Qu.:6.000	3rd Qu.:0.0000
Max. :7.000	Max. :7.000	Max. :1.0000

```
freedom %>%
  count(country) %>%
  arrange(n)
```

```
# A tibble: 193 x 2
```

country	n
---------	---

```

      <chr>           <int>
1 South Sudan         10
2 Montenegro          15
3 Serbia              18
4 Timor-Leste         22
5 Afghanistan         26
6 Albania             26
7 Algeria             26
8 Andorra             26
9 Angola              26
10 Antigua and Barbuda 26
# ... with 183 more rows

freedom %>% select(year) %>% table(useNA = "ifany")

.
1995 1996 1997 1998 1999 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010
  189   189   189   189   190   190   190   190   191   191   191   192   192   192   192   192
2011 2012 2013 2014 2015 2016 2017 2018 2019 2020
  193   193   193   193   193   193   193   193   193   193

freedom %>% select(status) %>% table(useNA = "ifany")

.
      F      NF      PF
2219 1257 1503

freedom %>% select(region_code) %>% table(useNA = "ifany")

.
      2      9      19   142   150
1388  364  910 1218 1099

freedom %>% select(region_name) %>% table(useNA = "ifany")

.
      Africa Americas      Asia      Europe      Oceania
      1388      910      1218      1099      364

freedom %>% select(is_ldc) %>% table(useNA = "ifany")

.
      0      1
3803 1176

freedom %>%
  distinct(country, region_name) %>% View()

```

We can see that it appears that not all countries have full data for all 26 years. There is data ranging from 1995 through 2020, with a total of 193 countries in the data set.

A precursory glance suggest that not all countries have the full year range of data. There are no unexpected unique values from our initial look at the data. Additionally, the `political_rights` and `civil_liberties` (renamed as `polical_rights` and `civil_liberties`) scores are within the expected ranges given the definitions provided by FreedomHouse.

Checking for NA values

```
freedom %>%
  summarize(across(everything(), ~ sum(is.na(.)))) %>%
  tidyr::pivot_longer(everything()) %>%
  arrange(desc(value)) %>%
  deframe()
```

country	year	civil_liberties	political_rights
0	0	0	0
status	region_code	region_name	is_ldc
0	0	0	0
country_code			
0			

NA counting methodology was taken from stackexchange

It appears that there are no NA values in this data set. This is reaffirming our findings from when we checked the summary and tables for the dataset earlier.

Recoding

Going forward, we will be using the `region_name` column in lieu of `Region_Code` for simplicity. Additionally, we will recode `is_ldc` into categorical values.

```
freedom = freedom %>%
  mutate(ldc = if_else(is_ldc == 1, "Yes", "No"),
         freedom_index = political_rights + civil_liberties)
```

```
freedom %>%
  select(is_ldc, ldc) %>% unique()
```

```
# A tibble: 2 x 2
  is_ldc ldc
  <dbl> <chr>
1     1 Yes
2     0 No
```

Grouping candidates

Let's examine the total number of unique entries per column to see good candidates for faceting or other categorization methods.

```
freedom %>%
  sapply(n_distinct)
```

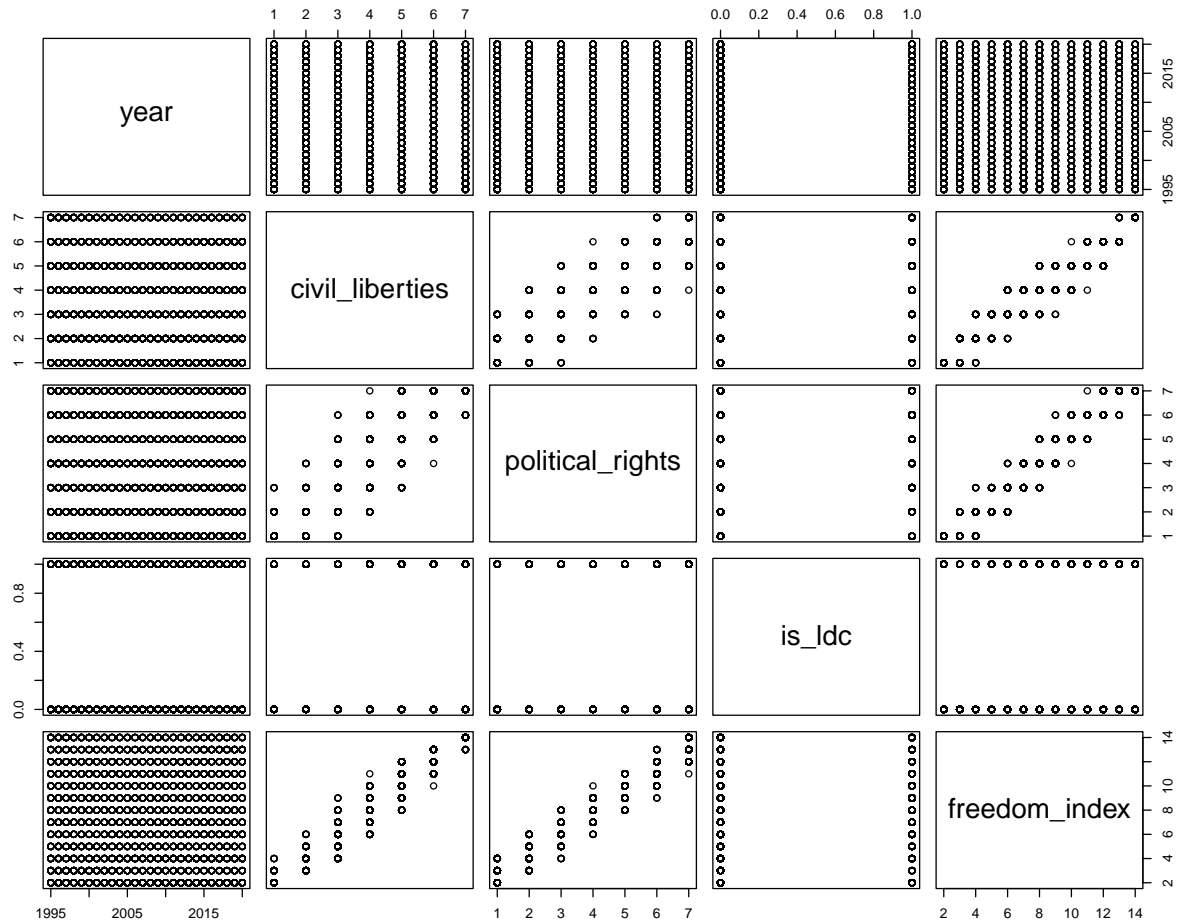
country	year	civil_liberties	political_rights
193	26	7	7
status	region_code	region_name	is_ldc
3	5	5	2
country_code	ldc	freedom_index	
193	2	13	

We can see that LDC, Status and `region_name` are all potential ways to cluster data.

Visual Exploration

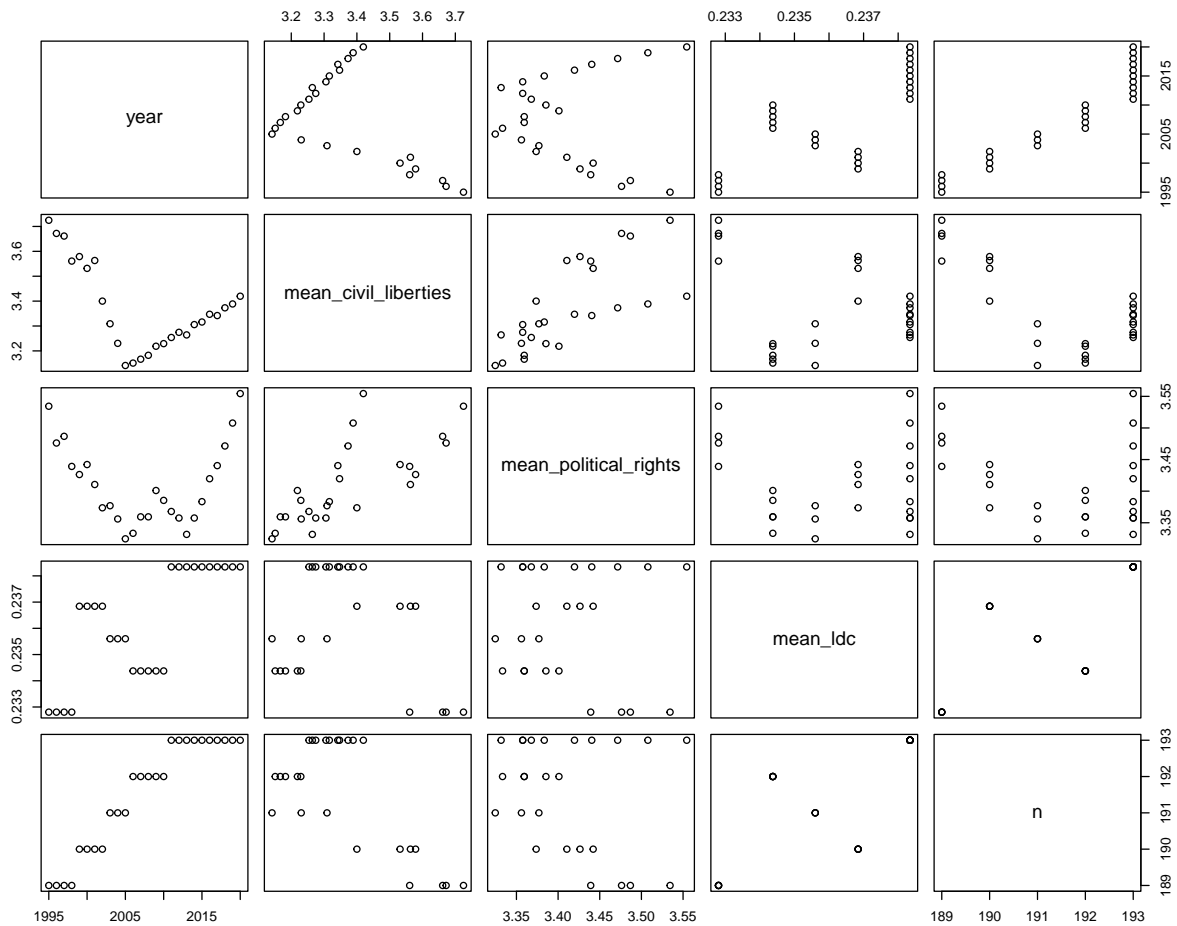
Pairwise plot for data - check raw data

```
freedom %>%
  select(year, civil_liberties, political_rights, is_ldc, freedom_index) %>%
  pairs()
```



Check grouped data

```
freedom %>%
  group_by(year) %>%
  summarize(mean_civil_liberties = mean(civil_liberties), mean_political_rights = mean(political_rights),
            mean_ldc = mean(is_ldc), n = n()) %>%
  select(year, mean_civil_liberties, mean_political_rights, mean_ldc, n) %>%
  pairs()
```



```
summarize_freedom <- function(tbl) {
  tbl %>%
    summarize(n_countries = n(),
              avg_civil_lib = mean(civil_liberties),
              avg_pol_rights = mean(political_rights),
              pct_free = mean(status == "F"),
              .groups = "drop") %>%
    arrange(desc(n_countries))
}
```

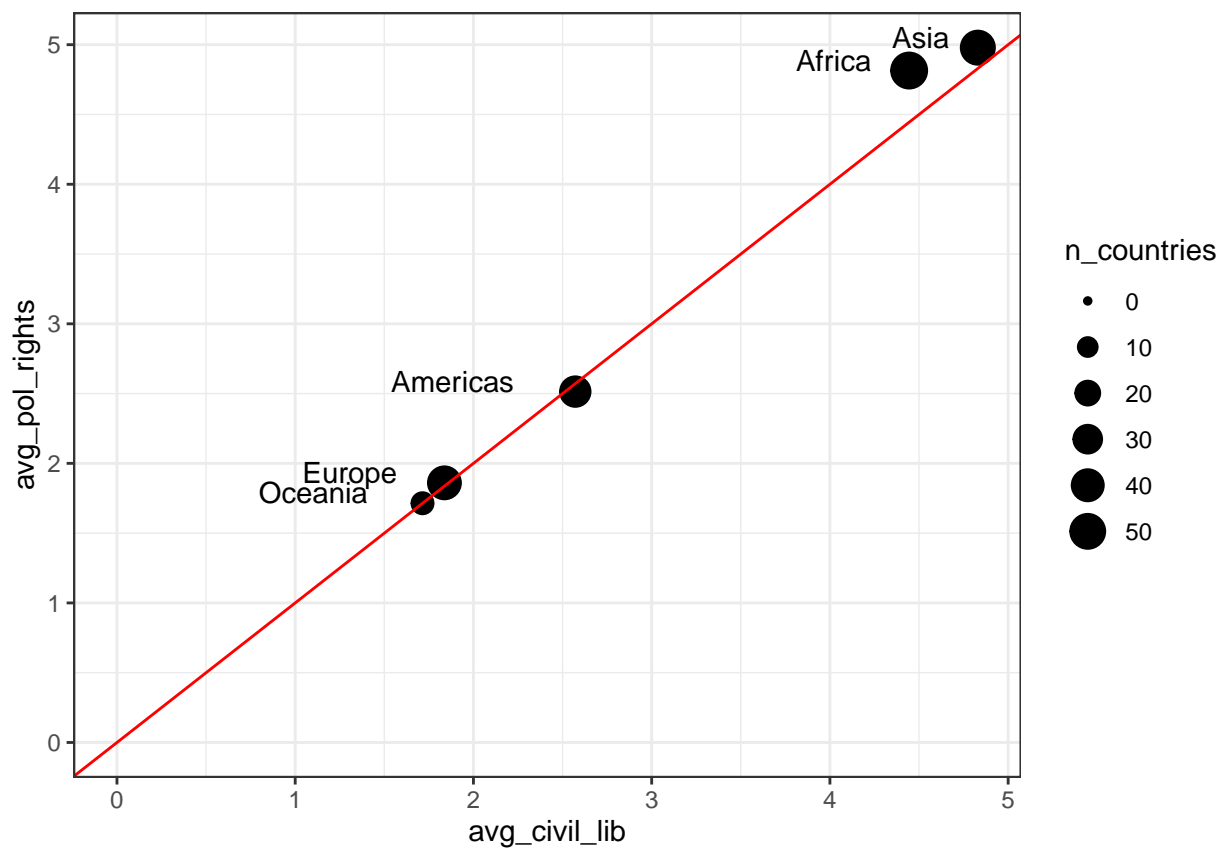
```
by_region <- freedom %>%
  filter(year == 2020) %>%
  group_by(region_name) %>%
  summarize(n_countries = n(),
            avg_civil_lib = mean(civil_liberties),
            avg_pol_rights = mean(political_rights),
            pct_free = mean(status == "F"))

by_region
```


```
# A tibble: 5 x 5
  region_name n_countries avg_civil_lib avg_pol_rights pct_free
```

	<chr>	<int>	<dbl>	<dbl>	<dbl>
1	Africa	54	4.44	4.81	0.167
2	Americas	35	2.57	2.51	0.6
3	Asia	47	4.83	4.98	0.128
4	Europe	43	1.84	1.86	0.767
5	Oceania	14	1.71	1.71	0.857

```
by_region %>% ggplot(aes(avg_civil_lib, avg_pol_rights)) +
  geom_point(aes(size = n_countries)) +
  geom_abline(color = "red") +
  geom_text(aes(label = region_name), vjust = 0, hjust = 1.5) +
  expand_limits(x = 0, y = 0, size = 0)
```



```
freedom %>%
  filter(year == 2020) %>%
  ggplot(aes(civil_liberties, political_rights)) +
  geom_jitter(height = 0.1, width = 0.1) +
  geom_abline(color = "red") +
  expand_limits(x = 0, y = 0, size = 0)
```

Developing Questions

While examining the data set, the following potential questions arose for investigation:

1. What is the distribution of freedom by Region in a single year?
2. How have the 5 of the best, and 5 of worst non LDC countries shifted in terms of freedom from 1995 to 2020?
3. How are proportions of Free, Partially Free, and Not Free countries shifting over time?
4. How are the distribution of Political Rights, and Civil Liberties fluctuating over time by region?

For the second question, we will need to set criteria for “worst” and “best”, and then identify these countries. We will use a combined `political_rights` + `civil_liberties` score as the gauge of measure. A score of 14 would represent “worst” and 2 would represent “best”. We will also exclude countries that do not have data for all years.

```
freedom %>%
  filter(year == 1995, ldc == "No") %>%
  select(country, region_name, ldc, freedom_index) %>%
  arrange(freedom_index) %>% head(20)
```

A tibble: 20 x 4

	country	region_name	ldc	freedom_index
	<chr>	<chr>	<chr>	<dbl>
1	Andorra	Europe	No	2
2	Australia	Oceania	No	2
3	Austria	Europe	No	2
4	Barbados	Americas	No	2
5	Belgium	Europe	No	2

6	Belize	Americas	No	2
7	Canada	Americas	No	2
8	Cyprus	Asia	No	2
9	Denmark	Europe	No	2
10	Dominica	Americas	No	2
11	Finland	Europe	No	2
12	Iceland	Europe	No	2
13	Ireland	Europe	No	2
14	Liechtenstein	Europe	No	2
15	Luxembourg	Europe	No	2
16	Malta	Europe	No	2
17	Marshall Islands	Oceania	No	2
18	Micronesia (Federated States of)	Oceania	No	2
19	Netherlands	Europe	No	2
20	New Zealand	Oceania	No	2

```
freedom %>%
  filter(year == 1995, ldc == "No") %>%
  select(country, region_name, ldc, freedom_index) %>%
  arrange(desc(freedom_index)) %>% head(20)
```

A tibble: 20 x 4

	country	region_name	ldc	freedom_index
	<chr>	<chr>	<chr>	<dbl>
1	China	Asia	No	14
2	Cuba	Americas	No	14
3	Equatorial Guinea	Africa	No	14
4	Iraq	Asia	No	14
5	Libya	Africa	No	14
6	Nigeria	Africa	No	14
7	Democratic People's Republic of Korea	Asia	No	14
8	Saudi Arabia	Asia	No	14
9	Syrian Arab Republic	Asia	No	14
10	Tajikistan	Asia	No	14
11	Turkmenistan	Asia	No	14
12	Uzbekistan	Asia	No	14
13	Viet Nam	Asia	No	14
14	Indonesia	Asia	No	13
15	Iran (Islamic Republic of)	Asia	No	13
16	Kenya	Africa	No	13
17	Qatar	Asia	No	13
18	Algeria	Africa	No	12
19	Azerbaijan	Asia	No	12
20	Bahrain	Asia	No	12

```
freedom %>%
  filter(year == 1995, ldc == "No", freedom_index == 2) %>%
  arrange(freedom_index) %>%
  select(country) %>%
  pull()
```

[1]	"Andorra"	"Australia"
[3]	"Austria"	"Barbados"
[5]	"Belgium"	"Belize"
[7]	"Canada"	"Cyprus"

```

[9] "Denmark" "Dominica"
[11] "Finland" "Iceland"
[13] "Ireland" "Liechtenstein"
[15] "Luxembourg" "Malta"
[17] "Marshall Islands" "Micronesia (Federated States of)"
[19] "Netherlands" "New Zealand"
[21] "Norway" "Portugal"
[23] "San Marino" "Sweden"
[25] "Switzerland" "United States of America"

```

```

freedom %>%
  filter(year == 1995, ldc == "No", freedom_index == 14) %>%
  arrange(freedom_index) %>%
  select(country) %>%
  pull()

```

```

[1] "China"
[2] "Cuba"
[3] "Equatorial Guinea"
[4] "Iraq"
[5] "Libya"
[6] "Nigeria"
[7] "Democratic People's Republic of Korea"
[8] "Saudi Arabia"
[9] "Syrian Arab Republic"
[10] "Tajikistan"
[11] "Turkmenistan"
[12] "Uzbekistan"
[13] "Viet Nam"

```

```

exclude_list <- freedom %>%
  group_by(country) %>%
  summarize(n = n()) %>%
  filter(n != 26) %>%
  select(country) %>% pull()
exclude_list

```

```

[1] "Montenegro" "Serbia" "South Sudan" "Timor-Leste"

```

A look at the data showed that there was not enough granularity in the scale to determine solely based off the data, thus we chose the following 5 countries based off interest and recent political events.

```

best_list <- c("Australia", "Canada", "Belgium", "Iceland", "United States of America")
worst_list <- c("China", "Iraq", "Nigeria", "Saudi Arabia", "Viet Nam")
choice_list <- c("Russian Federation", "Ukraine", "Qatar", "Afghanistan", "Rwanda")
best_list

```

```

[1] "Australia" "Canada"
[3] "Belgium" "Iceland"
[5] "United States of America"

```

```

worst_list

```

```

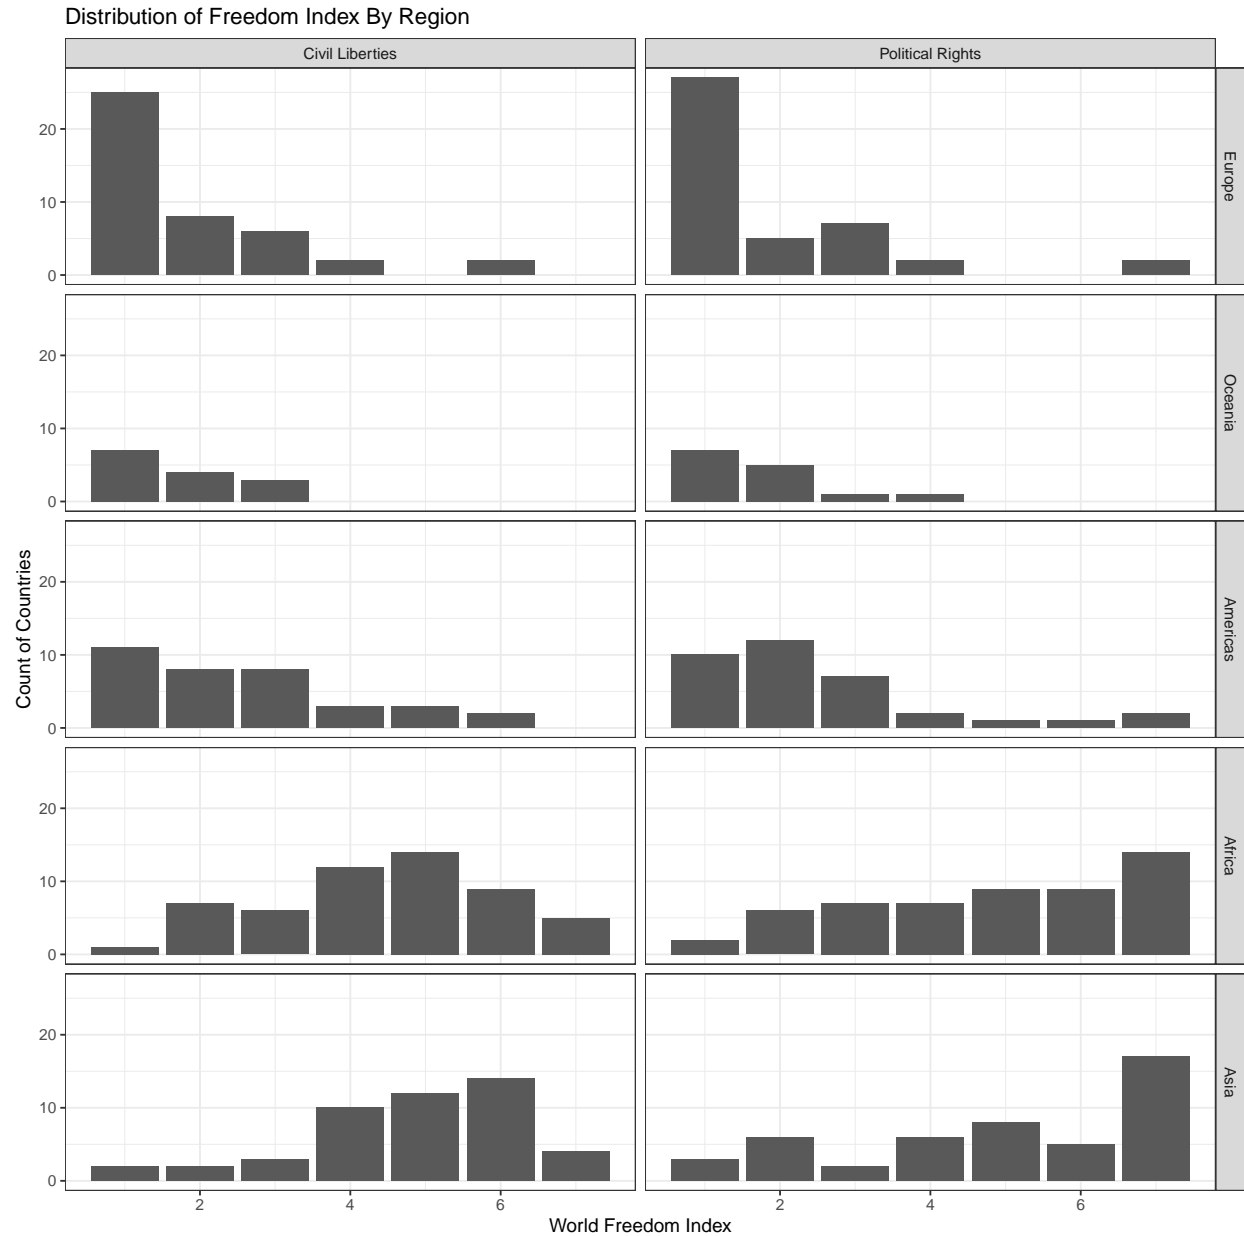
[1] "China" "Iraq" "Nigeria" "Saudi Arabia" "Viet Nam"

```

Data Visualizations

For the first question, we will look at the most recent year in the data set, 2020, and follow the analysis done by David Robinson:


```
freedom %>%
  filter(year == 2020) %>%
  gather(key = metric, value = value, civil_liberties, political_rights) %>%
  mutate(metric = str_to_title(str_replace_all(metric, "_", " ")),
         region_name = fct_reorder(region_name, value)) %>%
  count(region_name, metric, value) %>%
  ggplot(aes(value, n)) +
  geom_col() +
  facet_grid(region_name ~ metric) +
  scale_x_continuous(name = "World Freedom Index") +
  scale_y_continuous(name = "Count of Countries") +
  ggtitle("Distribution of Freedom Index By Region")
```




Trends for top and worst 5, and choice

```
freedom %>%
  filter(country %in% best_list) %>%
  pivot_longer(cols = c("political_rights", "civil_liberties"), names_to = "Label", values_to = "Score") +
  ggplot() +
  aes(x = year, y = Score, color = Label) +
  facet_wrap(
    vars(country),
    ncol = 1) +
  geom_line(size = 2) +
  scale_y_continuous(
    limit = c(0, 7)) +
```


```
scale_x_continuous(
  name = "Year") +
theme_classic()
```



```
freedom %>%
  filter(country %in% worst_list) %>%
  pivot_longer(cols = c("political_rights", "civil_liberties"), names_to = "Label", values_to = "Score") +
  ggplot() +
  aes(x = year, y = Score, color = Label) +
  facet_wrap(
    vars(country),
    ncol = 1) +
  geom_line(size = 2) +
  scale_y_continuous(
    limit = c(0, 7)) +
  scale_x_continuous(
    name = "Year") +
  theme_classic()
```



```
freedom %>%
  filter(country %in% choice_list) %>%
  pivot_longer(cols = c("political_rights", "civil_liberties"), names_to = "Label", values_to = "Score")
  ggplot() +
    aes(x = year, y = Score, color = Label) +
    facet_wrap(
      vars(country),
      ncol = 1) +
    geom_line(size = 2) +
    scale_y_continuous(
      limit = c(0, 7)) +
    scale_x_continuous(
      name = "Year") +
    theme_classic()
```



Unfortunately, these plots ended up poor visually. There was also limited movement in the metrics. We could speculate that maybe those with the highest and lowest freedom scores are most “stable” in regards to the operations of their current regime form. A notable exception among those countries scoring “high” in freedom is the U.S., in which this metric shows a decline in freedom starting in 2017.

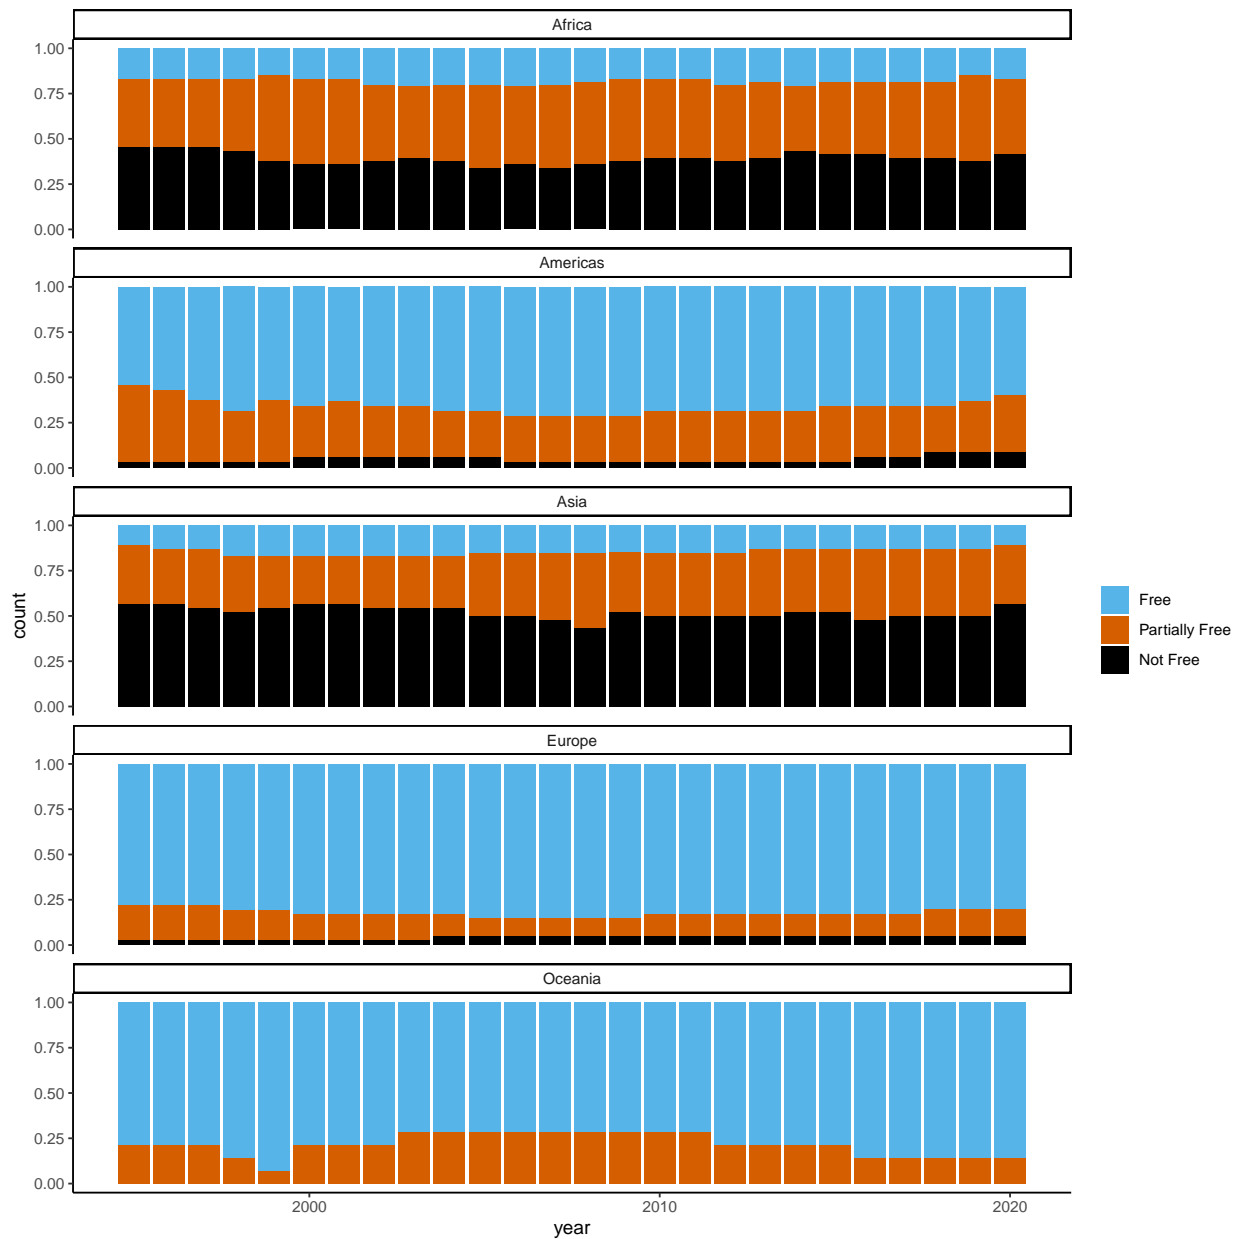
Overall trends in freedom over time

```
status_colors = c("#56B4E9", "#D55E00", "#000000")

freedom %>%
  filter(!country %in% exclude_list) %>%
  ggplot() +
  aes(x = year, fill = factor(status, levels = c("F", "PF", "NF")))) +
  geom_bar(position = "fill") +
  facet_wrap(
    vars(region_name),
    ncol = 1) +
  scale_fill_manual(
    name = "",
    values = status_colors,
    breaks = c("F", "PF", "NF"),
```




```
labels = c("Free", "Partially Free", "Not Free")) +
theme_classic()
```



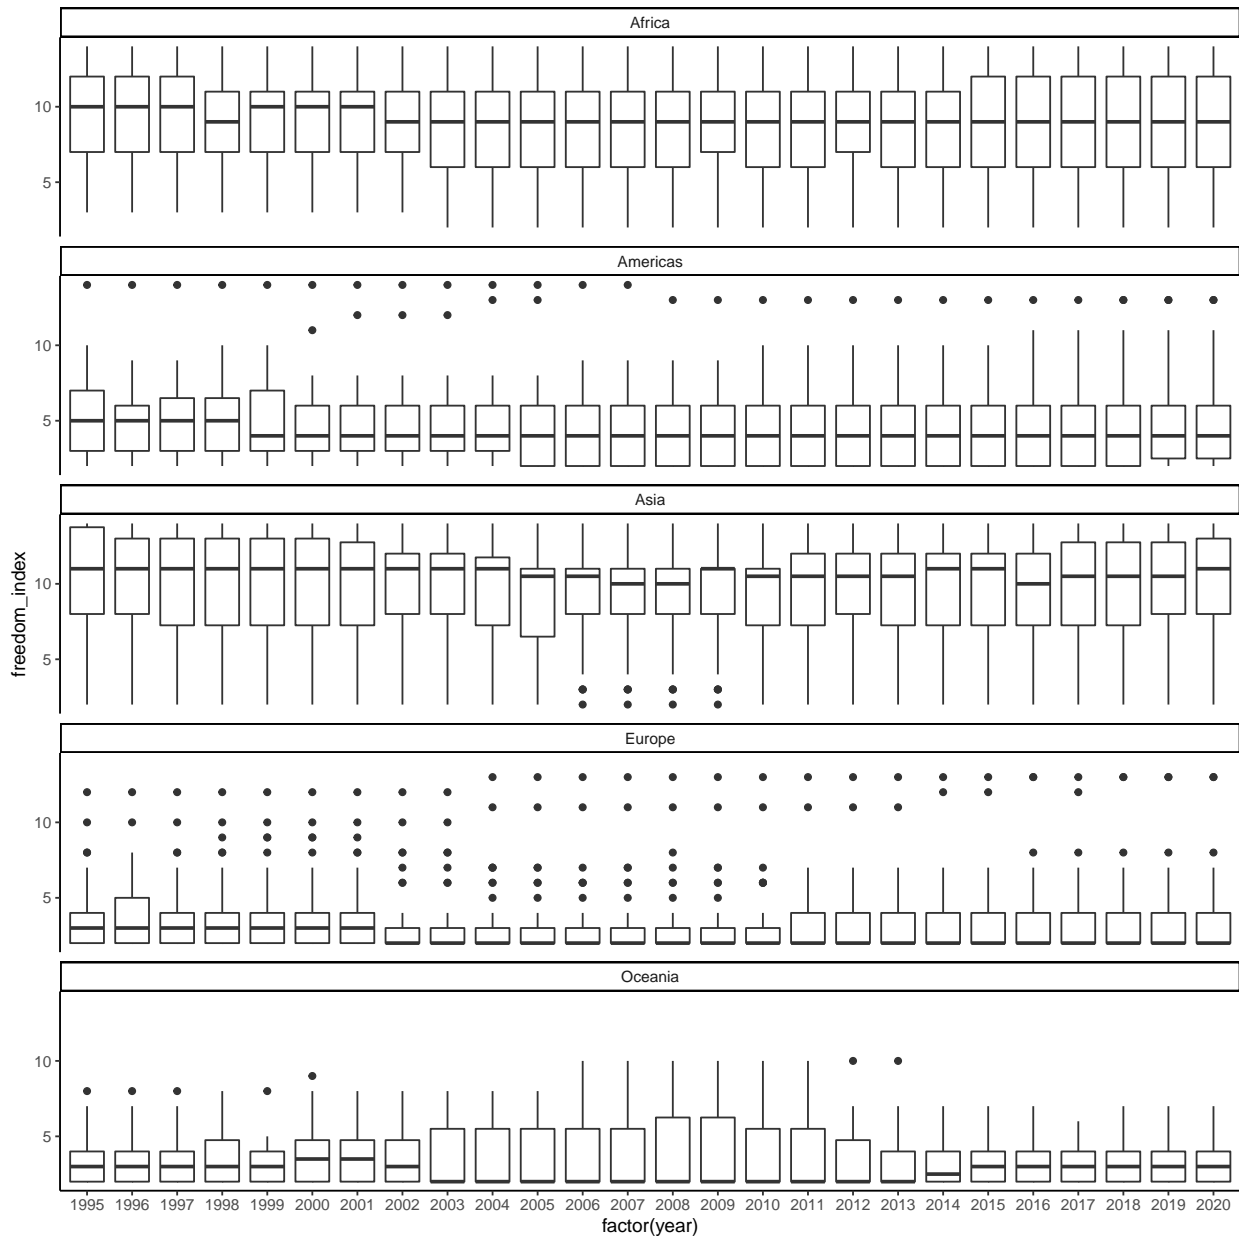
```
freedom %>%
  filter(!country %in% exclude_list) %>%
  ggplot() +
  aes(x = year, fill = factor(status, levels = c("F", "PF", "NF")))) +
  geom_bar(position = "fill") +
  scale_fill_manual(
    name = "",
    values = status_colors,
    breaks = c("F", "PF", "NF"),
    labels = c("Free", "Partially Free", "Not Free")) +
  ggtitle("Overall") +
```

```
theme_classic()
```



Overall trends in freedom over time by region

```
freedom %>%  
  filter(!country %in% exclude_list) %>%  
  ggplot() +  
  aes(x = factor(year), y = freedom_index) +  
  geom_boxplot() +  
  facet_wrap(  
    vars(region_name),  
    ncol = 1) +  
  theme_classic()
```



Box plots look quite messy,

We will re-evaluate with line charts as done in David Robinson's example:

```
freedom_gathered <- freedom %>%
  gather(key = metric, value = value, civil_liberties, political_rights) %>%
  mutate(metric = str_to_title(str_replace_all(metric, "_", " ")),
         region_name = fct_reorder(region_name, value))

overall <- freedom_gathered %>%
  group_by(year, metric) %>%
  summarize(avg_rating = mean(value))
```

`summarise()` has grouped output by 'year'. You can override using the `.groups` argument.

```

freedom_gathered %>%
  # mutate(region_name = "Overall") %>%
  # bind_rows(freedom_gathered) %>%
  group_by(year, region_name, metric) %>%
  summarize(avg_rating = mean(value)) %>%
  ggplot(aes(year, avg_rating)) +
  geom_line(size = 1, aes(color = region_name)) +
  geom_line(data = overall, size = 2) +
  facet_wrap(~ metric) +
  expand_limits(y = 1) +
  scale_x_continuous(name = "Year") +
  scale_y_reverse(name = "World Freedom Index", breaks = seq(1,7)) +
  scale_color_discrete(name = "Region", guide = guide_legend(reverse = TRUE)) +
  ggtitle("Change of Freedom Index over time by Region", subtitle = "Black Shows Overall")

```

`summarise()` has grouped output by 'year', 'region_name'. You can override using the `.groups` argument.

Change of Freedom Index over time by Region
Black Shows Overall



Additional Analysis joined with other data sets

We will perform some additional analysis by pulling in other data sets.

Worldbank Data

More information available at The World Bank.

```
#WDIsearch("gdp") %>% View()

gdp_percap <- WDI(indicator = "NY.GDP.PCAP.CD",
                  extra = TRUE,
                  start = 1995,
                  end = 2020) %>%
```

```

as_tibble()

# gdp_percap %>% View()

# freedom %>% distinct(country) %>% anti_join(gdp_percap, by = "country") %>% View()

# gdp_percap %>% filter(str_detect(country, "Egypt"))

freedom_join <- freedom %>%
  inner_join(gdp_percap,
    by = c(country_code = "iso2c", "year"),
    suffix = c("", "_wdi")) %>%
  mutate(income = fct_relevel(income, c("Low income", "Lower middle income", "Upper middle income")))

freedom_join %>%
  filter(year == 2020, income != "Not classified") %>%
  group_by(income) %>%
  summarize_freedom() %>%
  arrange(income)

# A tibble: 4 x 5
  income          n_countries avg_civil_lib avg_pol_rights pct_free
  <fct>              <int>      <dbl>         <dbl>      <dbl>
1 Low income          27        5.30          5.56        0
2 Lower middle income  54        4.02          4.15       0.222
3 Upper middle income  53        3.45          3.62       0.377
4 High income         58        1.91          1.95       0.845

freedom_join_gathered <- freedom_join %>%
  filter(income != "Not classified") %>%
  gather(key = metric, value = value, civil_liberties, political_rights) %>%
  mutate(metric = str_to_title(str_replace_all(metric, "_", " ")))

overall_join <- freedom_join_gathered %>%
  group_by(year, metric) %>%
  summarize(avg_rating = mean(value))

```

`summarise()` has grouped output by 'year'. You can override using the `.groups` argument.

```

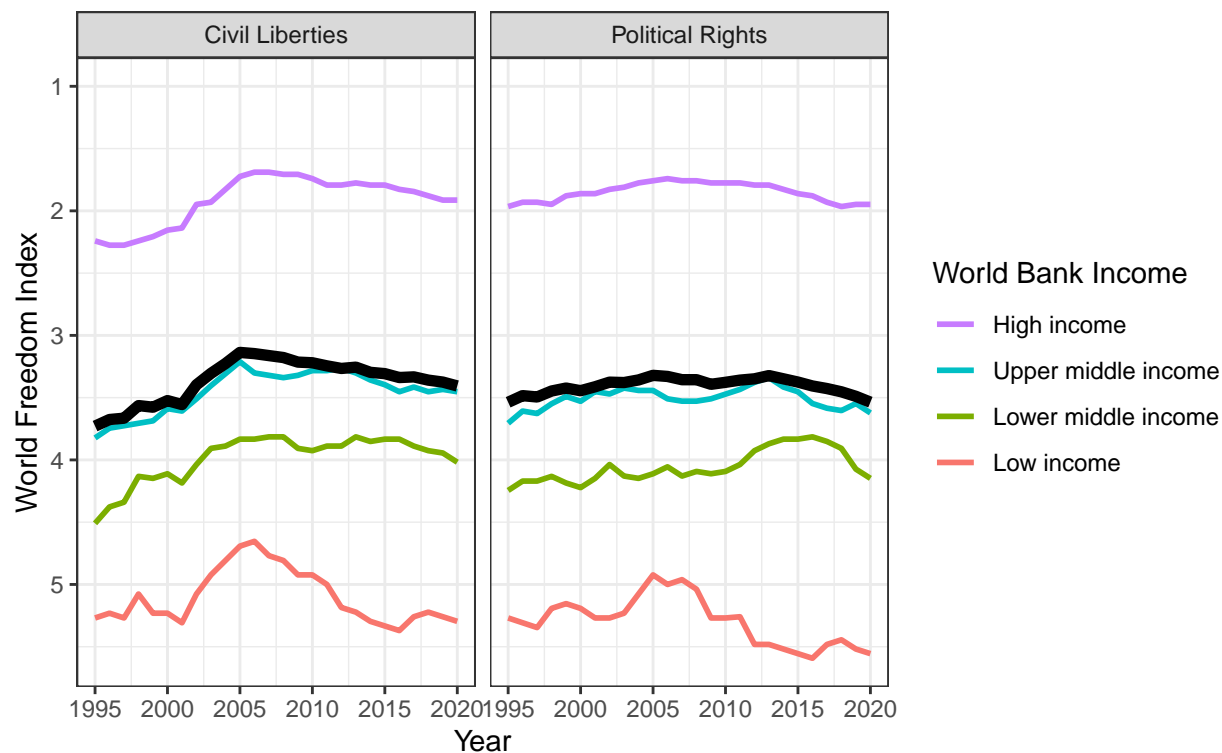
freedom_join_gathered %>%
  # mutate(region_name = "Overall") %>%
  # bind_rows(freedom_gathered) %>%
  group_by(year, income, metric) %>%
  summarize(avg_rating = mean(value)) %>%
  ggplot(aes(year, avg_rating)) +
  geom_line(size = 1, aes(color = income)) +
  geom_line(data = overall_join, size = 2) +
  facet_wrap(~ metric) +
  expand_limits(y = 1) +
  scale_x_continuous(name = "Year") +
  scale_y_reverse(name = "World Freedom Index", breaks = seq(1,7)) +
  scale_color_discrete(name = "World Bank Income", guide = guide_legend(reverse = TRUE)) +
  ggtitle("Change of Freedom Index over time by Income", subtitle = "Black Shows Overall")

```

`summarise()` has grouped output by 'year', 'income'. You can override using the `.groups` argument.

Change of Freedom Index over time by Income

Black Shows Overall



How do we find the biggest outliers?

freedom_join_gathered

```
# A tibble: 9,906 x 22
  country      year status region_code region_name is_ldc country_code ldc
  <chr>      <dbl> <chr>      <dbl> <chr>      <dbl> <chr>      <chr>
1 Afghanistan 1995 NF          142 Asia          1 AF          Yes
2 Afghanistan 1996 NF          142 Asia          1 AF          Yes
3 Afghanistan 1997 NF          142 Asia          1 AF          Yes
4 Afghanistan 1998 NF          142 Asia          1 AF          Yes
5 Afghanistan 1999 NF          142 Asia          1 AF          Yes
6 Afghanistan 2000 NF          142 Asia          1 AF          Yes
7 Afghanistan 2001 NF          142 Asia          1 AF          Yes
8 Afghanistan 2002 NF          142 Asia          1 AF          Yes
9 Afghanistan 2003 NF          142 Asia          1 AF          Yes
10 Afghanistan 2004 NF          142 Asia          1 AF          Yes
# ... with 9,896 more rows, and 14 more variables: freedom_index <dbl>,
#   country_wdi <chr>, NY.GDP.PCAP.CD <dbl>, status_wdi <chr>,
#   lastupdated <chr>, iso3c <chr>, region <chr>, capital <chr>,
#   longitude <chr>, latitude <chr>, income <fct>, lending <chr>, metric <chr>,
#   value <dbl>

civil_liberties_2020 <- freedom_join_gathered %>%
  filter(!is.na(NY.GDP.PCAP.CD),
```

```

    metric == "Civil Liberties",
    year == 2020)


lin_mod <- civil_liberties_2020 %>%
  lm(value ~ region_name + log10(NY.GDP.PCAP.CD), data = .)

lin_mod %>%
  augment(data = civil_liberties_2020) %>%
  select(country, region_name, NY.GDP.PCAP.CD, income, value, .fitted, .resid) %>%
  arrange(desc(abs(.resid))) %>%
  head(20) %>%
  ggplot(aes(.fitted, value)) +
  geom_point() +
  geom_abline(color = "red") +
  geom_text_repel(aes(label = country)) +
  scale_x_reverse(name = "Model Civil Liberty") +
  scale_y_reverse(name = "Actual Civil Liberty") +
  expand_limits(x = 1, y = 1) +
  ggtitle("Outliers for Civil Liberties", subtitle = "Lower Number is Better")

```

Outliers for Civil Liberties

Lower Number is Better



Map Visual

```

freedom_2020 <- freedom_join_gathered %>%
  filter(year == 2020,
    metric == "Civil Liberties")

```

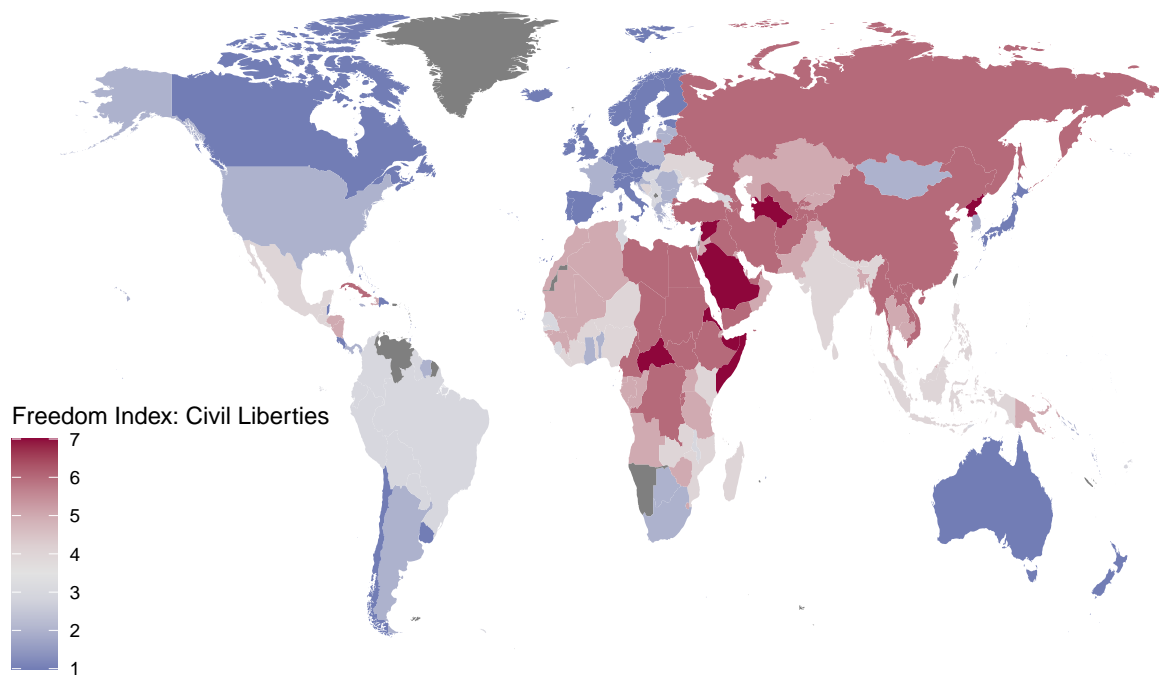


```
# maps::iso3166

world_map_freedom_2020 <- map_data("world") %>%
  as_tibble() %>%
  regex_left_join(maps::iso3166, c(region = "mapname")) %>%
  left_join(freedom_2020 %>% select(-region), by = c(a2 = "country_code")) %>%
  filter(region != "Antarctica")

world_map_freedom_2020 %>%
  ggplot(aes(long, lat, group = group)) +
  geom_polygon(aes(fill = value)) +
  coord_fixed(1.5) +
  scale_fill_continuous_diverging(palette = "Blue-Red", mid = 3.5, name = "Freedom Index: Civil Liberties")
ggtitle("Civil Liberties Index", subtitle = "In 2020") +
ggthemes::theme_map()
```

Civil Liberties Index
In 2020



Conclusion

This data set was fairly clean and had a large amount of pre-processing completed on it. My impression is that it would have been better for Country-level analysis if the data set had supplied the raw indicator scores which composed the civil_liberties and political_rights metrics. As can be seen in the “Top 5” and “Worst 5” bar charts, the level of granularity provided gives minimal year by year changes. Region level analysis has more apparent trends when examining both by region and changes over time. Lastly, by bringing in economic indicators, we can examine the correlation between freedom and income.